

Week 7: Mixed Models

MATH-516 Applied Statistics

Tomas Masak

Feb 20th 2023

Section 1

Course Organization (Update)

- **Week 1:** Intro
 - Project 1: Snow Data
- Week 2: Linear Models - Practical Recap
- **Week 3:** Logistic Regression
 - Project 2: Online Shopping Data
- Week 4: Generalized Linear Models
- **Week 5:** Poisson Regression
 - Project 3: Premier League Data
- Week 6: more on GLMs
- **Week 7:** Linear Mixed Models
 - Project 4: U.S. Presidential Elections
- Free Week: Easter Holidays
- Week 8: Linear Mixed and Multilevel Models

Content (cont.)

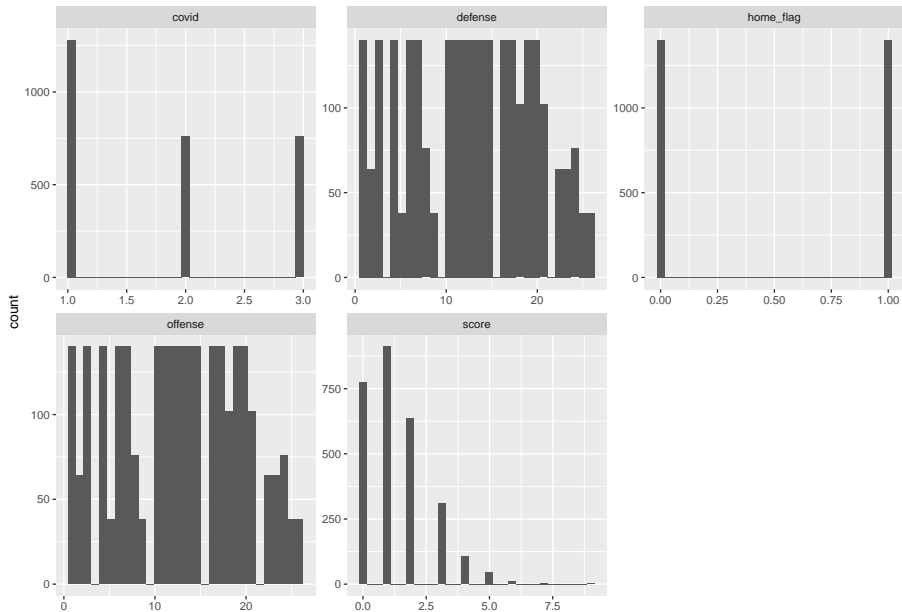
- **Week 9:** Time Series
 - Project 5: Global Warming
- Week 10: Time Series Regression
- **Week 11:** Functional Data Analysis
 - Project 7: First Wave of Covid in the US
- Week 12: Functional PCA
- **Week 13:** Statistical Consulting
- Week 14: **Oral Exam**
 - discussing your submitted projects

Evaluation: remains the same as announced on Week 1, with Statistical Consulting replacing Project 7 (active participation + writing up a suggested solution to a presented problem). Work on Projects 6 and 7 can be combined (to be considered a single submission) and will not be examined.

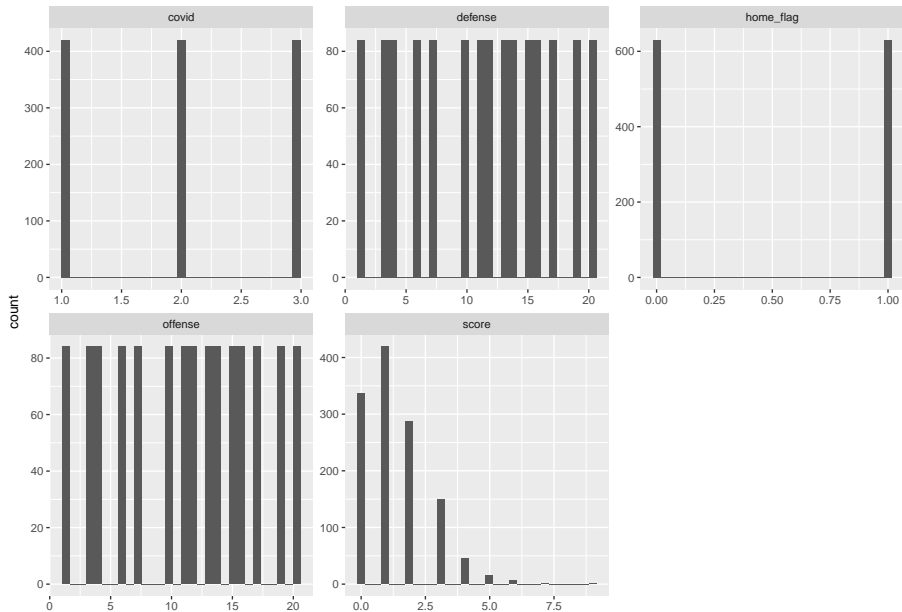
Section 2

Project 3: Premier League

Data Visualized



Balanced Data



Models (Balanced Data)

```
m <- glm(score~covid*home_flag+defense+offense, data=subDat, family="poisson")
msub <- glm(score~home_flag+defense+offense, data=subDat, family="poisson")
anova(m,msub,test="LRT")
```

```
## Analysis of Deviance Table
##
## Model 1: score ~ covid * home_flag + defense + offense
## Model 2: score ~ home_flag + defense + offense
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      1226      1352.8
## 2      1230      1357.5 -4   -4.7254  0.3167
```


Models (Full Data)

```
m <- glm(score~covid*(home_flag+defense+offense), data=Data, family="poisson")
msub <- glm(score~home_flag+defense+offense, data=Data, family="poisson")
anova(m,msub,test="LRT")
```

```
## Analysis of Deviance Table
##
## Model 1: score ~ covid * (home_flag + defense + offense)
## Model 2: score ~ home_flag + defense + offense
##   Resid. Df Resid. Dev  Df Deviance Pr(>Chi)
## 1         2674      2952.6
## 2         2748      3057.7 -74   -105.14   0.0101 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
library(car)
Anova(m,type=2)
```

```
## Analysis of Deviance Table (Type II tests)
##
## Response: score
##           LR Chisq Df Pr(>Chisq)
## covid           0.91  2  0.633917
## home_flag       16.99  1 3.761e-05 ***
## defense        184.35 25 < 2.2e-16 ***
## offense        345.44 25 < 2.2e-16 ***
## covid:home_flag   5.77  2  0.055895 .
## covid:defense     60.04 35  0.005292 **
## covid:offense     36.35 35  0.405718
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Models (Full Data)

```
minter <- glm(score~covid*(home_flag+defense)+offense, data=Data, family="poisson")
msub <- glm(score~home_flag+covid*defense+offense, data=Data, family="poisson")
anova(m,msub,test="LRT")
```

```
## Analysis of Deviance Table
##
## Model 1: score ~ covid * (home_flag + defense + offense)
## Model 2: score ~ home_flag + covid * defense + offense
##   Resid. Df Resid. Dev  Df Deviance Pr(>Chi)
## 1      2674      2952.6
## 2      2711      2994.7 -37   -42.102   0.2596
anova(minter,msub,test="LRT")
```

```
## Analysis of Deviance Table
##
## Model 1: score ~ covid * (home_flag + defense) + offense
## Model 2: score ~ home_flag + covid * defense + offense
##   Resid. Df Resid. Dev  Df Deviance Pr(>Chi)
## 1      2709      2988.9
## 2      2711      2994.7 -2    -5.7544  0.05629 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
sum(is.na(coefficients(m)))

## [1] 30
```

- we are on the edge of significance with a model that has too few observations to rely on asymptotics and to estimate all the parameters

Section 3

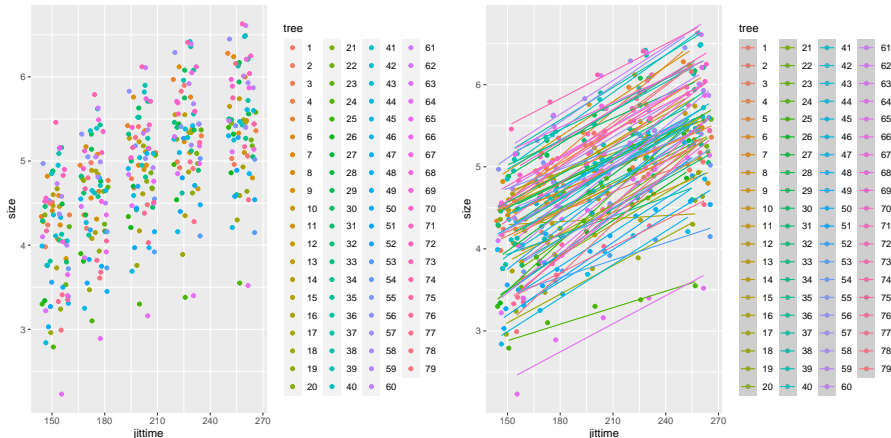
Example: Tree Growth

Data

- log-size ($\log\text{-height} + 2\log\text{-diameter}$) of 79 (Sitka spruce) trees measured repeatedly in about 1-month intervals
 - each tree measured 5-times
 - 54 trees grown in ozone-enriched environment ($\text{treat}=1$) and 25 were control

```
##      size time tree treat
## 1 4.51   152     1 ozone
## 2 4.98   174     1 ozone
## 3 5.41   201     1 ozone
## 4 5.90   227     1 ozone
## 5 6.15   258     1 ozone
## 6 4.24   152     2 ozone
```

Data Displayed



Right: individual line for every tree corresponding to model

$y \sim \text{tree} * \text{time}$

Models

```
m1 <- lm(size~(time+I(time^2))*tree,data=Sitka)
m0 <- lm(size~(time+I(time^2)), data=Sitka)
anova(m0,m1)
```

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	391	157.107				
2	237	6.267	154	150.84	37.043	< 2.2e-16 ***

- m_0 allows for a separate curve for the two treatment groups
- m_1 allows for a separate curve for every tree (so m_0 is a submodel of m_1)

Problems:

- m_1 cannot be simplified to m_0 , but the effect of interest (treat) cannot be fitted without this simplification, because every single tree is either treatment or control
 - also what if we had low number of observations for some trees and couldn't afford to fit m_1 ?
- but assumptions of m_0 are clearly violated

Diagnostics





Figure: Residuals grouped by tree.

Section 4

Linear Mixed Models

Definition

- regressions with a large no. of coefficients some of which are themselves being modelled
- extend the linear model

$$Y = \mathbf{X}\beta + \epsilon, \quad \epsilon \sim \mathcal{N}_N(0, \sigma^2 \mathbf{I}_N)$$

to the **linear mixed model**

$$Y = \mathbf{X}\beta + \mathbf{Z}b + \epsilon, \quad b \sim \mathcal{N}_q(0, \mathbf{C}), \quad \epsilon \sim \mathcal{N}_n(0, \sigma^2 \mathbf{I}_n)$$

- \mathbf{X} and \mathbf{Z} are known design matrices
- $\beta \in \mathbb{R}^p$ are fixed (non-random) parameters (effects)
- $b \in \mathbb{R}^q$ are random effects with mean 0 and covariance matrix \mathbf{C}
 - independent of ϵ
- parameters: β , \mathbf{C} and σ^2

Fitting the Model (ML method)

- the linear mixed model has its log-likelihood $\ell(\beta, \mathbf{C}, \sigma^2)$
- if we knew \mathbf{C} , we could rewrite the model to

$$Y = \mathbf{X}\beta + e, \quad e \sim \mathcal{N}_n(0, \sigma^2 \mathbf{W}), \quad \mathbf{W} = \mathbf{I}_n + \frac{1}{\sigma^2} \mathbf{Z}\mathbf{C}\mathbf{Z}^\top$$

then the solution $\hat{\beta}_{\mathbf{C}}$ and $\hat{\sigma}_{\mathbf{C}}^2$ would be given explicitly by weighted least squares

- imagine a reparametrization “if we knew \mathbf{C}/σ^2 ” instead
- consider the profile log-likelihood for \mathbf{C} : $\ell_p(\mathbf{C}) = \ell(\hat{\beta}_{\mathbf{C}}, \mathbf{C}, \hat{\sigma}_{\mathbf{C}}^2)$

Algorithm: starting from an initial $\mathbf{C}^{(0)}$ alternate until convergence for $l = 1, 2, \dots$ between

- 1 calculation of $\hat{\beta}_{\mathbf{C}^{(l-1)}}$ and $\hat{\sigma}_{\mathbf{C}^{(l-1)}}^2$ by weighted least squares
- 2 updating $\mathbf{C}^{(l)} = \arg \min_{\mathbf{C}} \ell_p(\mathbf{C})$ by Newton's method (itself iterative)
 - since Newton works well given good starting values, one rather runs EM algorithm for a while (treating b as unobserved data) before switching to this scheme

Fitting the Model (REML method)

- start by integrating out β from the log-likelihood
 - it actually has a closed form and is equivalent to working with likelihood for $\mathbf{A}Y$ such that $\mathbb{E}\mathbf{A}Y = 0$, i.e. it is just a simple transformation of the problem
- use iterative solver
 - closed form for σ^2
 - inner iteration for \mathbf{C}
- finally obtain β as with the ML method

ML vs. REML:

- REML is often preferred (and set as default) since it can lead to unbiased variance estimators (which ML never does)
- but REML depends on parametrization of the fixed effects
 - if one wants to compare models with different \mathbf{X} using likelihood criteria, ML needs to be used!
- ML and REML are asymptotically equivalent

Example: Tree Growth

Rewrite the model in a form that shows the grouping:

$$Y_k = \mathbf{X}_k \beta + \mathbf{Z}_k b_k + \epsilon, \quad b_k \sim \mathcal{N}_q(0, \mathbf{C}), \quad \epsilon_k \sim \mathcal{N}_{N_k}(0, \sigma^2 \mathbf{I}_{N_k})$$

- b_i and ϵ_i are i.i.d. for $k = 1, \dots, K$
- $\{b_i\}_{k=1, \dots, K} \perp \{\epsilon_k\}_{k=1, \dots, K}$

Specifically, in the tree growth example m1 can be replaced by:

- K is the no. of trees
- $N_k = 5$ (for all k) is the number of measurements per tree
-

$$Y_k = \begin{pmatrix} Y_{k1} \\ \vdots \\ Y_{k5} \end{pmatrix} \quad \mathbf{X}_k = \begin{pmatrix} 1 & t_{k1} & t_{k1}^2 \\ \vdots & \vdots & \vdots \\ 1 & t_{k5} & t_{k5}^2 \end{pmatrix} = \mathbf{Z}_k$$

-

$$\beta = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{pmatrix} \quad b_k = \begin{pmatrix} b_{k0} \\ b_{k1} \\ b_{k2} \end{pmatrix} \quad \mathbf{C} = \begin{pmatrix} \sigma_0^2 & \sigma_{01} & \sigma_{02} \\ \sigma_{01} & \sigma_1^2 & \sigma_{12} \\ \sigma_{02} & \sigma_{12} & \sigma_2^2 \end{pmatrix}$$

Example & Predictors of Random Effects

Fixed effect model: $\mathbb{E}[Y_{ki} | X_{ki} = t_{ki}] = \beta_{k0} + \beta_{k1}t_{k1} + \beta_{k2}t_{ki}^2$

- or something similar depending on the parametrization of `tree`, e.g. with `contr.sum` it would be for all but the last tree:

$$\mathbb{E}[Y_{ki} | X_{ki} = t_{ki}] = (\beta_0 + \beta_{k0}) + (\beta_1 + \beta_{k1})t_{ki} + (\beta_2 + \beta_{k2})t_{ki}^2$$

Mixed effect model:

$$\mathbb{E}[Y_{ki} | X_{ki} = t_{ki}, b_k] = (\beta_0 + b_{k0}) + (\beta_1 + b_{k1})t_{ki} + (\beta_2 + b_{k2})t_{ki}^2$$

- $\text{cov}(Y_{ki}, Y_{kj}) = \text{cov}(b_{k0} + b_{k1}t_{ki} + b_{k2}t_{ki}^2, b_{k0} + b_{k1}t_{kj} + b_{k2}t_{kj}^2) \neq 0$

In general (no need to read too much into the formula):

$$\hat{b}_k = \mathbf{CZ}_k^\top (\mathbf{Z}_k \mathbf{CZ}_k^\top + \sigma^2 \mathbf{I}_{N_k})^{-1} (Y_k - \mathbf{X}_k^\top \hat{\beta})$$

- called *predictors* since these are not parameters but random variables, but they are also sort of *shrinkage estimators*, because
- vaguely: \hat{b}_{kl} is somewhere between 0 and the $\hat{\beta}_{kl}$ in the `contr.sum` parametrization above

Example: Tree Growth

Consider a simpler model, where only the intercept is random:

$$\mathbb{E}[Y_{ki} \mid X_{ki} = t_{ki}, b_k] = (\beta_0 + b_k) + \beta_{k1}t_{ki}$$

and the corresponding fixed-effect-only model `y ~ tree+time`.




Figure: Tree Growth data and lines by fixed-effect-only model (dashed) and random intercept model (solid).

Uncertainty Quantification

Let $\theta \in \mathbb{R}^r$ denote the vector of parameters determining \mathbf{C} .

Theorem. Under validity of the model above and the MLE regularity conditions, we have for $K \rightarrow \infty$:

① estimators $\hat{\beta}, \hat{\theta}, \hat{\sigma}^2$ are consistent,

②
$$\sqrt{K} \begin{pmatrix} \hat{\beta} - \beta \\ \hat{\theta} - \theta \\ \hat{\sigma}^2 - \sigma^2 \end{pmatrix} \rightarrow \mathcal{N}_{p+r+1}(0, \mathbf{J}^{-1}), \text{ where } \mathbf{J} = \begin{pmatrix} \mathbf{J}_{\beta} & 0 & 0 \\ 0 & \mathbf{J}_{\theta} & \mathbf{J}_{\theta, \sigma^2} \\ 0 & \mathbf{J}_{\theta, \sigma^2} & \mathbf{J}_{\sigma^2} \end{pmatrix}$$

is the Fisher information matrix,

③ when $\hat{\ell}$ denotes the maximized log-likelihood of the model and $\hat{\ell}_0$ denotes the maximized log-likelihood of a submodel then

$$2[\hat{\ell} - \hat{\ell}_0] \rightarrow \chi_m^2,$$

where m is the difference in the no. of parameters between the model and the submodel.

Testing for Model Components

- testing fixed effects, i.e. $H_0 : \beta_{p-m+1} = \dots = \beta_p = 0$ against $H_1 : \neg H_0$
 - can be done via LRT due to point 3. of the previous theorem
 - ML needs to be used instead of REML
 - however, p-values tend to be too small, sometimes overstating importance of some effects
- testing random effects, i.e. $H_0 : \theta_{p-m+1} = \dots = \theta_p = 0$
 - usually cannot be done using the previous theorem, because MLE regularity assumptions are typically not met
 - one of the components of θ is typically the variance for one of the components of b_i 's, which lies on the edge of the parameter space
 - it can still be shown that the LR statistic still follows χ^2 -distribution, but with a smaller no. of degrees of freedom
 - p-values of the test the LRT from point 3. of the previous theorem are too big, sometimes understating importance of some effects

While solutions based on theory exist, a simpler road for us is the parametric bootstrap.

Example: Tree Growth

```
library(lme4)
# standardize time, otherwise convergence issues
Sitka <- Sitka %>% mutate(time=(time-mean(time))/sd(time))
mm1 <- lmer(size~treat*(time+I(time^2)) + (time+I(time^2)|tree),
            data=Sitka,REML=F)
mm0 <- lmer(size~time+I(time^2) + (time+I(time^2)|tree),
            data=Sitka,REML=F)
anova(mm0,mm1)

## Data: Sitka
## Models:
## mm0: size ~ time + I(time^2) + (time + I(time^2) | tree)
## mm1: size ~ treat * (time + I(time^2)) + (time + I(time^2) | tree)
##      npar      AIC      BIC logLik deviance Chisq Df Pr(>Chisq)
## mm0     10 -111.88 -72.089 65.939  -131.88
## mm1     13 -118.50 -66.772 72.249  -144.50 12.62  3   0.005535 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- REML=F because the fixed-effect structure differs between the two models
- treatment seems significant, but... let's do the bootstrap

Example: Tree Growth

Bootstrap still rejects, although the p-value is doubled:

```
lrstat <- as.numeric(2*(logLik(mm1)-logLik(mm0)))
lrstats <- rep(0,1000)
for(i in 1:1000){
  set.seed(517*i)
  if(i %% 10 ==0) print(i)
  newDat <- Sitka
  newDat$size <- unlist(simulate(mm0))
  bnull <- lmer(size~time+I(time^2) + (time+I(time^2)|tree),
               data=newDat,REML=F)
  balt <- lmer(size~treat*(time+I(time^2)) + (time+I(time^2)|tree),
               data=newDat,REML=F)
  lrstats[i] <- as.numeric(2*(logLik(balt)-logLik(bnull)))
}
mean(lrstats > lrstat)
```

```
[1] 0.011
```







So we can finally conclude that ozone treatment matters :)

Similar to the standard linear model, with the additional

- check of normality for the random effects

```
plot(fitted(mm1),resid(mm1)) # the only thing `plot(mm1)` gives
qqnorm(resid(mm1),main="") # a bit heavy tails
plot(cooks.distance(mm1),type="h") # the old ROTs not useful here
abline(h=3*mean(cooks.distance(mm1)),col="red",lty=2) # another ROT
qqnorm(ranef(mm1)$tree[,1], main="QQ for Intercept")
qqnorm(ranef(mm1)$tree[,2], main="QQ for Linear Term")
qqnorm(ranef(mm1)$tree[,3], main="QQ for Quadratic Term")
# this is only checking marginals; multivariate GoF tests are tricky
```

Diagnostics



summary(mm1)

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
tree	(Intercept)	0.356490	0.59707	
	time	0.013509	0.11623	0.04
	I(time^2)	0.002593	0.05092	0.12 -0.71
Residual		0.008096	0.08998	

Number of obs: 395, groups: tree, 79

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	5.093296	0.120072	42.419
treatozone	-0.201347	0.145231	-1.386
time	0.546397	0.024638	22.177
I(time^2)	-0.108450	0.014041	-7.724
treatozone:time	-0.079045	0.029800	-2.652
treatozone:I(time^2)	-0.009835	0.016984	-0.579

- the random quadratic term is very significant (LR test's p-value 10^{-8} even though understate), but of a low variance and high correlation ... problems

Section 5

Multilevel Models

Grouping

There can be more than one type of grouping:

$$Y_{kl} = \mathbf{X}_{kl}\beta + \mathbf{Z}_{k,l}^{(1)}b_k + \mathbf{Z}_{l,k}^{(2)}b_l + \epsilon_{kl}$$

- Y_{kl} in $\mathbb{R}^{N_{kl}}$ is the vector of individual observations belonging to
 - k -th level of grouping (1)
 - l -th level of grouping (2)
- e.g. if in the tree growth example we needed to group not only by tree but also by time
 - all i -th observations per tree could be correlated
 - e.g. if a different person would measure tree size for different times i , but the same person for a single i

Or groupings can be nested:

$$Y_{kl} = \mathbf{X}_{kl}\beta + \mathbf{Z}_{k,l}b_k + \mathbf{Z}_{kl}b_{kl} + \epsilon_{kl}$$

- these are called multilevel models:
 - $k = 1, \dots, K$ is the first-level grouping (e.g. school)
 - $l = 1, \dots, L$ is the second-level grouping (e.g. class)
 - $Y_{kl} \in \mathbb{R}^{N_{kl}}$ is the vector of individual observations (e.g. students)

R syntax (lme4 package)

formula	meaning
(1group)	random group intercept
$(x \text{group}) = (1+x \text{group})$	random slope of x within group with correlated intercept
$(0+x \text{group}) = (-1+x \text{group})$	random slope of x within group: no variation in intercept
$(1 \text{group}) + (0+x \text{group})$	uncorrelated random intercept and random slope within group
$(1 \text{site/block}) = (1 \text{site})+(1 \text{site:block})$	intercept varying among sites and among blocks within sites (nested random effects)
$\text{site}+(1 \text{site:block})$	<i>fixed</i> effect of sites plus random variation in intercept among blocks within sites
$(x \text{site/block}) = (x \text{site})+(x \text{site:block}) = (1 + x \text{site})+(1+x \text{site:block})$	slope and intercept varying among sites and among blocks within sites
$(x1 \text{site})+(x2 \text{block})$	two different effects, varying at different levels
$x*\text{site}+(x \text{site:block})$	fixed effect variation of slope and intercept varying among sites and random variation of slope and intercept among blocks within sites
$(1 \text{group1})+(1 \text{group2})$	intercept varying among crossed random effects (e.g. site, year)

source: [link](#)

- the linear mixed model generalizes to the GLMM in the same way that the standard linear model generalizes to the GLM
- GLMMs are not as useful, because β 's can only be interpreted as the population-average effects if the link g is linear:

$$\mathbb{E}[Y_n | X_n] = \mathbb{E}g^{-1}(X_n^\top \beta + Z_n^\top b_n)$$

- cannot go inside g^{-1} with the expectation unless it is linear, i.e. cannot get rid of b_n unless g^{-1} is linear
- only Gaussian models have the linear link as the canonical link
- if one does not use a canonical link, issues may arise
 - numerical convergence issues
 - calculated estimators are not guaranteed to be MLEs \Rightarrow theory does not work, etc.

Section 6

Project 4

Data

- data on U.S. Presidential elections between 1948 and 1992
- two-party vote, i.e. any other candidate than the Democratic and the Republican parties disregarded
- the data are already heavily pre-processed with many derived variables
 - details below
 - contains polls, etc.

In the 90s, there was a general belief that U.S. presidential elections are easy to understand using linear models and the kind of data we have.

Goal of the project: Endorse or refute the general belief by:

- start by exploring the data and building a linear model
- consider whether observations in a given a year and/or in a given region within year are correlated
- motivate using mixed models and build one

More detailed description of data and tasks can be found in `./Misc/Projec-4_assignment`.