

# Week 10: Forecasting and Time Series Regression

## MATH-516 Applied Statistics

Tomas Masak

May 1st 2023

## Section 1

# Forecasting with SARIMA

# Truth about Forecasting



**Forecasting time series is like driving a car using the rear mirror only.**


- though arguably more useful as a skill (the field is huge!)

Depending on who you ask SARIMA models are either:

- fairly old and simple, but surprisingly powerful
- overcomplicated and easily outperformed by vanilla methods
  - e.g. fitting a regression model and extrapolating

# Air Passengers Revisited

```
library(forecast) # needed on every line below
# fit <- auto.arima(AirPassengers)
fit <- Arima(AirPassengers, order=c(2,1,1), seasonal=c(0,1,0)) # capital A
preds <- forecast(fit, 12*10, c(0.5,0.95)) # 2-year ahead forecast
autoplot(AirPassengers) + autolayer(preds)
```



# Explanation

$$\Phi_{[P]}(B^s)\phi_{[P]}(B)(1 - B^s)^D(1 - B)^dX_t = \Theta_{[Q]}(B^s)\theta_{[Q]}(B)Z_t$$

How are the point forecasts above calculated? In a recursive way:

- 1 Express  $X_t$  from the SARIMA equation above.
- 2 Replace  $t$  by  $N + h$  in the equation.
- 3 Replace future observations (up to  $N + h - 1$ ) by their forecasts, future errors by zeros, and past errors by the respective residuals. (these are BLUPs given the past)

Under certain assumptions:

- $Y_t = (1 - B^s)^D(1 - B)^dX_t$  must be
  - stationary
  - *invertible*, i.e. polynomial  $\Theta_{[Q]}(x^s)\theta_{[Q]}(x)$  has no roots inside the unit circle, i.e.  $Z_t$  can be expressed as a linear combination of past  $X$ 's
    - typically we also require *causality*, i.e. polynomial  $\Phi_{[P]}(x^s)\phi_{[P]}(x)$  has no roots inside the unit circle, i.e.  $X_t$  can be expressed as a linear combination of past  $Z$ 's

the procedure above results in the best linear unbiased prediction. (under Gaussianity, it is the best prediction in the mean square sense)

# ARIMA(1,1,1) Example

Say that we fitted an ARIMA(1,1,1) model

$$(1 - \hat{\phi}B)(1 - B)X_t = (1 + \hat{\theta}B)Z_t$$

The forecasting procedure above consists of the following steps:

- ①  $X_t = X_{t-1} + \hat{\phi}X_{t-1} - \hat{\phi}X_{t-2} + Z_t + \hat{\theta}Z_{t-1}$
- ②  $X_{N+1} = X_N + \hat{\phi}X_N - \hat{\phi}X_{N-1} + Z_{N+1} + \hat{\theta}Z_N$
- ③  $\hat{X}_{N+1} = X_N + \hat{\phi}X_N - \hat{\phi}X_{N-1} + 0 + \hat{\theta}\hat{Z}_N$

This was a one-step ahead forecast. To do two steps ahead, replace  $N$  by  $N + 1$  in step 2 and repeat (replacing  $X_{N+1}$  by  $\hat{X}_{N+1}$  that is now available).

# Prediction Intervals

- constructed under Gaussian assumptions
  - i.e. as  $\hat{X}_{N+h} \pm q_{1-\alpha/2} \hat{\sigma}_h$ , where  $q_{1-\alpha/2}$  is a Gaussian quantile and  $\hat{\sigma}_h$  is the standard error of the prediction, calculated with estimators plugged in (details omitted), hence
- only reflects uncertainty about future realizations of  $Z$ 's!
- including uncertainty about parameter estimation possible via Monte Carlo simulation
  - i.e. parametric bootstrap, drawing parameter values for every simulation from their respective asymptotic distribution (was not discussed, but exists)
  - also possible to do non-parametric bootstrap instead to protect ourselves against non-Gaussian  $Z$ 's (set `bootstrap=T` in `forecast()`)
- including uncertainty about model selection requires a more careful simulation study
  - a relatively easy but somewhat arbitrary option is to simulate from an unnecessarily large model and perform automatic model selection for every simulation run

## Section 2


# Time Series Regression



# Air Passengers Revisited

- say the trend in Air Passengers data set is linear (after log) and we want a confidence interval on the slope.

```
lmData <- data.frame(y = log(AirPassengers),  
                     month = as.factor(rep(1:12,length(AirPassengers)/12)),  
                     t = 1:length(AirPassengers))  
lmfit <- lm(y~., data=lmData)  
plot(log(AirPassengers))  
points(seq(1949, 1961, length=144),fitted(lmfit), type="l", col="blue")
```



# Air Passengers Revisited

```
summary(lmfit)$coefficients
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)  4.726780368 0.0188935382 250.1797341 1.966475e-177
## month2      -0.022054823 0.0242108707  -0.9109471  3.639964e-01
## month3       0.108172299 0.0242117524   4.4677600  1.691103e-05
## month4       0.076903445 0.0242132220   3.1760930  1.861594e-03
## month5       0.074530803 0.0242152792   3.0778420  2.539909e-03
## month6       0.196677004 0.0242179239   8.1211339  2.980730e-13
## month7       0.300619331 0.0242211560  12.4114362  7.210080e-24
## month8       0.291324492 0.0242249751  12.0257912  6.615459e-23
## month9       0.146689890 0.0242293811   6.0542153  1.392768e-08
## month10      0.008531649 0.0242343735   0.3520474  7.253684e-01
## month11      -0.135186061 0.0242399521  -5.5769937  1.343178e-07
## month12      -0.021321065 0.0242461164  -0.8793600  3.808163e-01
## t            0.010068805 0.0001193001  84.3989506  4.561377e-116
```

```
confint(lmfit)["t",]
```

```
##          2.5 %      97.5 %
## 0.009832801 0.010304809
```

- Problem: data are not i.i.d.
- going back to the proof of the Gauss-Markov theorem, we would realize
  - the estimates are still unbiased (does not mean much)
  - they are no longer “best” (among linear unbiased estimators)
  - the standard errors, and hence also CIs, are wrong (thinking there is more df)

# Generalized Least Squares

- standard least squares (regression) model:
  - $Y = \mathbf{X}\beta + \epsilon$  with  $\epsilon \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$
- generalized least squares model:
  - $Y = \mathbf{X}\beta + \epsilon$  with  $\epsilon \sim \mathcal{N}(\mathbf{0}, \sigma^2 \Sigma)$
- if we knew the dependency structure  $\Sigma$ , we could change the variables:
  - $\Sigma^{-1/2}Y = \Sigma^{-1/2}\mathbf{X}\beta + \Sigma^{-1/2}\epsilon$
  - $\tilde{Y} = \tilde{X}\beta + \tilde{\epsilon} \dots$  standard least squares work here (and they correspond to Gaussian MLE)
- if we knew the dependency structure up to a fixed no. of parameters  $\alpha \in \mathbb{R}^q$ , we could do the Gaussian MLE jointly for  $\beta, \sigma^2, \alpha$ 
  - this is exactly the case when the residuals follow an ARMA model
  - one could also iterate between the change of variables and fitting ARMA to the residuals
    - this is called Cochrane-Orcutt method, and it provably converges to the MLE, which is usually obtained numerically (`glts()` from the `nlme` package in R)

# Air Passengers Revisited

- lets say that AR(1) is a good model for the residuals above
  - this is what we can decide doing time series analysis like last week

```
library(nlme)
R_struct <- corARMA(form=~t, p=1) # AR(1) correlation structure
glsfit <- gls(y~t+month, data=lmData, corr=R_struct, method="ML")
confint(lmfit)["t",]
```

```
##          2.5 %          97.5 %
## 0.009832801 0.010304809
```

```
confint(glsfit)["t",] # wider than above
```



```
##          2.5 %          97.5 %
## 0.009364262 0.010624771
```

# Air Passengers Revisited

- fitted values are typically very similar, and so are the residuals
  - if they are not, we should re-think our dependency structure

```
plot(as.numeric(lmData$y),type="l")  
points(fitted(lmfit),type="l", col="red")  
points(fitted(glsfit),type="l", col="blue", lty=2)
```

```
plot(resid(lmfit),type="l", col="red")  
points(as.numeric(resid(glsfit)),type="l", col="blue", lty=2)
```



# Hypothesis Testing

Maximum likelihood theory (potentially with nuisance parameters) works here:

- let  $Y \sim \mathcal{N}_N(\mathbf{X}\beta, \sigma^2\Sigma)$ , where  $\Sigma \in \mathbb{R}^{N \times N}$  depends on a parameter vector  $\alpha \in \mathbb{R}^r$  (of a fixed sized as  $N \rightarrow \infty$ )
- let  $\theta \in \mathbb{R}^p$  is a vector containing all the parameters, i.e.  $\beta, \sigma^2, \alpha$ , such that the hypothesis concerns the first  $q$  entries of  $\theta$ , namely:
- let  $\Theta_1 \times \Theta_c$ , where  $\Theta_1 \subset \mathbb{R}^q$  and  $\Theta_c \in \mathbb{R}^{p-q}$  is the parameter space, and  $\Theta_0 \subset \Theta_1$ 
  - $H_0 : \theta \in \Theta_0 \times \Theta_c$
  - $H_1 : \theta \in \Theta_1 \times \Theta_c$
- then under  $H_0$  the likelihood ratio statistics approaches the  $\chi^2$ -distribution:

$$LR_N = 2[\ell_N(\hat{\theta}_1) - \ell_N(\hat{\theta}_0)] \sim \chi_q^2$$

- $\hat{\theta}_0 = \arg \max_{\theta \in \Theta_0 \times \Theta_c} \ell_N(\theta)$
- $\hat{\theta}_1 = \arg \max_{\theta \in \Theta_1 \times \Theta_c} \ell_N(\theta)$

# Air Passengers Revisited

- `anova()` in R works
  - use `method="ML"` is used instead of "REML" when changing the mean structure

```
library(car)
# glsfit <- gls(y~t+month, data=lmData, corr=R_struct, method="ML")
Anova(glsfit, type=2)
```

```
## Analysis of Deviance Table (Type II tests)
```

```
##
```

```
## Response: y
```

```
##           Df  Chisq Pr(>Chisq)
```

```
## t           1 966.02  < 2.2e-16 ***
```

```
## month      11 911.92  < 2.2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# or half-manually:
```

```
subfit_1 <- gls(y~month, data=lmData, corr=R_struct, method="ML")
```

```
subfit_2 <- gls(y~t, data=lmData, corr=R_struct, method="ML")
```

# Air Passengers Revisited

```
anova(subfit_1, glsfit)
```

```
##           Model df          AIC          BIC   logLik   Test L.Ratio p-value
## subfit_1      1 14 -498.1899 -456.6125 263.0950
## glsfit        2 15 -526.1789 -481.6317 278.0895 1 vs 2  29.989  <.0001
```

```
anova(subfit_2, glsfit)
```

```
##           Model df          AIC          BIC   logLik   Test L.Ratio p-value
## subfit_2      1  4 -251.3892 -239.5100 129.6946
## glsfit        2 15 -526.1789 -481.6317 278.0895 1 vs 2 296.7897  <.0001
```

```
# or entirely manually:
```

```
1-pchisq(2*(glsfit$logLik - subfit_1$logLik),
         length(coef(glsfit)) - length(coef(subfit_1)))
```

```
## [1] 4.345036e-08
```


```
1-pchisq(2*(glsfit$logLik - subfit_2$logLik),
         length(coef(glsfit)) - length(coef(subfit_2)))
```

```
## [1] 0
```



# Beaver Body Temperature

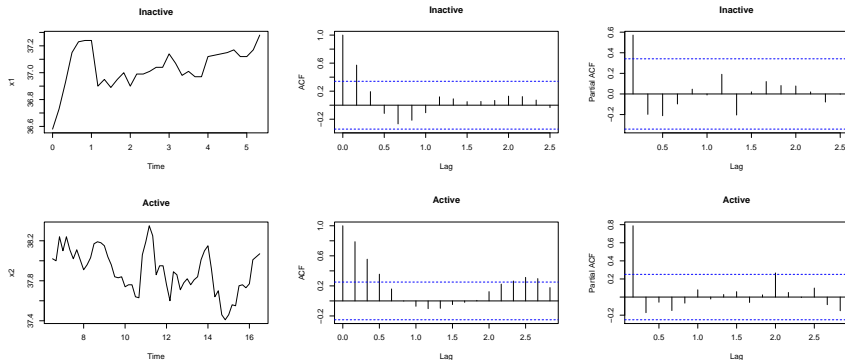
```
data(beavers)
beaver2 <- beaver2[,c(3,4)] # drop time
x <- ts(beaver2, start=0,frequency=6)
plot(x, main="", xlab="time [h]")
```



**Question:** What is the effect of activity on the body temperature?

# Beaver Body Temperature

```
x1 <- window(x, start=0, end=max(which(x[,2]==0))/6-1)[,1]
x2 <- window(x, start=min(which(x[,2]==1))/6)[,1]
plot(x1,main="Inactive"); acf(x1, main="Inactive"); pacf(x1, main="Inactive")
plot(x2, main="Active"); acf(x2, main="Active"); pacf(x2, main="Active")
```



Lags have no meaning here, let's take AR(1) as the model

# Beaver Body Temperature

```
beaver2$time <- seq_along(beaver2$temp)
R_struct <- corARMA(form=~time, p=1) # AR(1) correlation structure
( glsfit <- gls(temp~activ, data=beaver2, corr=R_struct, method="ML") )

## Generalized least squares fit by maximum likelihood
##   Model: temp ~ activ
##   Data: beaver2
##   Log-likelihood: 66.77523
##
## Coefficients:
## (Intercept)      activ
##  37.1919453    0.6141776
##
## Correlation Structure: AR(1)
##   Formula: ~time
##   Parameter estimate(s):
##       Phi
## 0.8731771
## Degrees of freedom: 100 total; 98 residual
## Residual standard error: 0.2527856
```

# Beaver Body Temperature (Alternatively)

```
arima(x[,1], c(1,0,0), xreg=x[,2])
```

```
##
```

```
## Call:
```

```
## arima(x = x[, 1], order = c(1, 0, 0), xreg = x[, 2])
```

```
##
```

```
## Coefficients:
```

```
##          ar1  intercept  x[, 2]
```

```
##          0.8733     37.1920  0.6139
```

```
## s.e.    0.0684      0.1187  0.1381
```

```
##
```

```
## sigma^2 estimated as 0.01518:  log likelihood = 66.78,  aic = -125.55
```

- standard errors are slightly different with `arima()`, since calculated empirically, while `glm()` uses asymptotic formula
- in both cases, activity increases body temperature of the beaver by about 0.6 degrees
- we could test significance of this value as above (but we see from s.e. that clearly significant)